

[Hadoop](#)[Blog](#)[Try the Sandbox!](#)[Community & Support](#)[Stay Connected](#)

Overview Lab 1: Getting started with BigInsights

[Events](#) [Videos](#)[saracco](#)Published on August 22, 2014 / *Updated on March 28, 2015*

2

[All Documentation](#)[Installing IBM BigIntegrate](#)

▼ [BigInsights and IBM Open Platform](#)

▼ [Getting started](#)

- [Skills Roadmap: New to Hadoop](#)
- [Skills Roadmap: Exploring BigInsights](#)
- [Skills Roadmap: Advanced Users](#)
- [Partners](#)

▼ [Tutorials](#)

▶ [Big SQL on Hadoop tutorial \(4.1\)](#)

- [Big SQL web tooling \(Data Server Manager\)](#)

▶ [HBase introductory tutorial \(4.0\)](#)

▼ [Overview tutorial \(3.0\)](#)

- [Overview Lab 1: Getting started with Hadoop and BigInsights](#)

- [Overview Lab 2: Issuing basic Hadoop commands](#)
- [Overview Lab 3: Using the web console to explore and administer your cluster](#)
- [Overview Lab 4: Analyzing social media data with BigSheets](#)
- [Overview Lab 5: Introducing Big SQL](#)
- [FAQs](#)
- [PAM support in Ambari for IOP 4.3](#)
- ▶ [Open source](#)
- ▶ [BigInsights value add](#)
- ▶ [Integration](#)

Lab 1: Overview and getting started

In this hands-on lab, you'll learn how to work with Big Data using Apache Hadoop and Info platform. In particular, you'll learn the basics of working with the Hadoop Distributed File System (HDFS) in a Hadoop-based environment using the BigInsights Web console. After launching a sample Hadoop-based environment using the BigInsights Web console, you'll explore a sample data set in a sophisticated scenario involving social media data. In doing so, you'll learn how to use a spreadsheet-style tool to analyze data about the global coverage of a popular brand without writing any code. Finally, you'll learn how to query data managed by BigInsights through IBM's Big SQL technology. Indeed, you'll have a chance to explore data over data in HDFS, including data derived from a relational data warehouse.

Ready to get started?

After completing this hands-on lab, you'll be able to:

- Work directly with Apache Hadoop through file system commands
- Inspect and administer your cluster through the BigInsights Web Console
- Explore big data using a spreadsheet-style tool
- Use Big SQL to create tables and issue complex queries

Allow 2 ½ – 3 hours to complete all labs associated with *Exploring Hadoop and BigInsights* in this section, plus additional time to download and unzip the required VMware image.

This lab was developed by Cynthia M. Saracco, IBM Silicon Valley Lab. Please post questions

technologies it describes to the forum on Hadoop Dev at <https://developer.ibm.com/answe>

1.1.About your environment

This lab was developed for the InfoSphere BigInsights 3.0 Quick Start Edition VMware image single-node cluster VMware image from this site: <http://www-01.ibm.com/software/data/i>
[start/downloads.html](http://www-01.ibm.com/software/data/i)

The VMware image is set up in the following manner:

	User	Password
VM Image root account	root	password
VM Image lab user account	biadmin	biadmin
BigInsights Administrator	biadmin	biadmin
Big SQL Administrator	bigsql	bigsql
Lab user	biadmin	biadmin

Property	Value
Host name	bivm.ibm.com
BigInsights Web Console URL	http://bivm.ibm.com:8080
Big SQL database name	bigsql
Big SQL port number	51000

**About the screen captures, sample code, and environment confi**

Screen captures in this lab depict examples and results what you see when you complete the exercises. In add examples may need to be customized to match your en example, you may need to alter directory path informat information.

1.2. Getting started

To get started with the lab exercises, you need to install and launch the VMware image as v

__1. If necessary, obtain a copy of the BigInsights 3.0 Quick Start Edition VMware image fr 01.ibm.com/software/data/infosphere/biginsights/quick-start/downloads.html). Use the i

__2. Follow the instructions provided to decompress (unzip) the file and install the image (file with additional information.

__3. If necessary, install VMware player or other required software to run VMware images. the BigInsights VMware image.

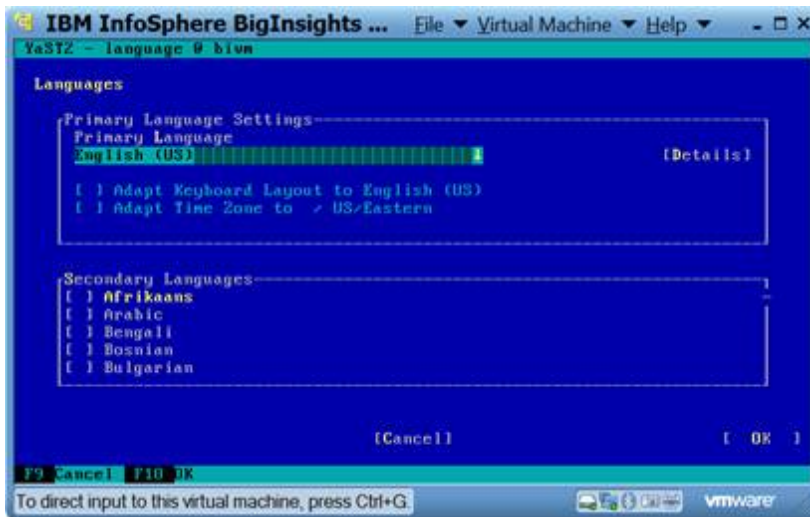
__4. Launch the VMware image. When logging in for the first time, use the root ID (with a p to configure your environment, accept the licensing agreement, and enter the passwords f and biadmin/biadmin) when prompted. **This is a one-time only requirement.**

```
IBM InfoSphere BigInsights ... File Virtual Machine Help
Mount CIFS File Systems                               unused
Loading console font lat9w-16.psfu -n trivial 60:loadable done
Loading keymap assuming iso-8859-15 euro               done
Loading /usr/share/kbd/keymaps/i386/german/us.map.gz   done
Loading compose table latin1.add                      done
Start Unicode mode                                     done
Starting irqbalance                                   done
Starting mcelog...                                     done
Setting up (remotefs) network interfaces:              done
Setting up service (remotefs) network . . . . . done
Starting SSH daemon                                   done
Starting Name Service Cache Daemon                    done
Starting mail service (Postfix)                       done
Starting CRON daemon                                  done
Starting smartd                                        unused
Master Resource Control: runlevel 3 has been           reached
Skipped services in runlevel 3:                        nfs smbfs splash smartd


Welcome to SUSE Linux Enterprise Server 11 SP2 (x86_64) - Kernel 3.0.101-0.5-default (tty1).

bium login: root
Password:

To direct input to this virtual machine, press Ctrl+G.
```



__5. When the one-time configuration process is completed, you will be presented with a :
with a password of biadmin.



__6. Verify that your screen appears similar to this:



__7. Click **Start BigInsights** to start all required services. (Alternatively, you can open a terminal window and run the command: `$BIGINSIGHTS_HOME/bin/start-all.sh`)



Wait until the operation completes. This may take several minutes, depending on your machine.

__8. Verify that all required BigInsights services are up and running. From a terminal window, run the command: `$BIGINSIGHTS_HOME/bin/status.sh`.

__9. Inspect the results, a subset of which are shown below. Verify that, at a minimum, the following services are successfully started: hdm, zookeeper, hadoop, catalog, hive, bigsql, oozie, console, and https.

```

[INFO] Progress - Status hbase
[INFO] @bivm.ibm.com - hbase-master(active) started, pid 8407
[INFO] @bivm.ibm.com - hbase-regionserver started, pid 8573
[INFO] Deployer - hbase service started
[INFO] Progress - 45%
[INFO] Progress - Status hive
[INFO] @bivm.ibm.com - hive-web-interface started, pid 9746
[INFO] @bivm.ibm.com - hive-server2 started, pid 9997
[INFO] Progress - 55%
[INFO] Progress - Status bigsql
[INFO] @bivm.ibm.com - bigsql-server started, pid 10622
[INFO] @bivm.ibm.com - head-node started, pid 15033
[INFO] @bivm.ibm.com - work-node started, pid 15028
[INFO] @bivm.ibm.com - scheduler started, pid 16026
[INFO] Progress - 65%
[INFO] Progress - Status oozie
[INFO] @bivm.ibm.com - oozie started, pid 17831
[INFO] Progress - 75%
[INFO] Progress - Status console
[INFO] @bivm.ibm.com - console started, pid 18138
[INFO] Progress - 82%
[INFO] Progress - Status https
[INFO] @bivm.ibm.com - https-server started, pid 20332
[INFO] Progress - 91%
[INFO] Progress - Status alert
[INFO] @bivm.ibm.com - alert started, pid 21034
[INFO] Progress - 100%
[INFO] DeployManager - Status: SUCCEEDED components: [hdm, zookeeper, hadoop, catalog, hbase, hive, bigsql, oozie, console, https, alert]; Consumes : 13506ms
biadm@bivm:/opt/ibm/biginsights/bin>

```

Now you're ready to start working with big data!

To find the other tutorials in this series, go to [Overview tutorial](#).



If have any questions or need help getting your environment up and runni visit Hadoop Dev (<https://developer.ibm.com/hadoop/>) and review the produ documentation or post a message to the forum.

You cannot proceed with subsequent lab exercises until you've logged int VMware image and launched the necessary BigInsights services.

TAGS OVERVIEW, SQL, BIG SQL, BIGINSIGHTS VALUE ADD, WEB CONSOLE, BIGSHEETS, MAPREDUCE, CYNTHIA M. S STARTED, NEW TO BIGINSIGHTS

saracco

2 comments on"Overview Lab 1: Getting started with Hadc

Justin Murray * • October 28, 2016

A question: does IBM have other images of BigInsights that come in VMware – such that I coul multiple virtual machines from an example template? Thank you.

[Reply](#)

Anwar • July 10, 2018

the download link is not working, please update it accordingly

[Reply](#)

Join The Discussion

Your email address will not be published. Required fields are marked *

Enter your comments...

Name *

Email *

Website

[Contact](#) [Privacy](#) [Terms of use](#) [Accessibility](#) [Report Abuse](#) [Feedback](#) [Cook](#)